

HW 5 By Sayem Lincoln

PID - A54207835

1. (70 points) Complete the linear regression functions (least squares and ridge regression) in the class. Use the code to conduct the following experiments:

(a) (10 points) Randomly generate 30 data points from the sine function, where each data point $(x; y)$ has the form:

$x = [x_0; x_1; x_2; \dots; x_{10}]; x \sim [0; 2\pi]$

$y = \sin(x) + \epsilon; \epsilon \sim N(0; 0.3)$

(b) (10 points) Plot the data points along with the sine function.

(c) (20 points) Randomly split the dataset (you can use the function provided in the Regression code) and use 30% of the data points for training and the rest for testing. Apply ridge regression using $\lambda = [1e-10; 1e-5; 1e-2; 1e-1; 1; 10; 100; 1000]$. Plot the training and testing performance.

(d) (30 points) Implement s-fold cross validation. Use $k = 4$ to choose the optimal λ from the set λ above.

Solution

```
def generate_sin_data(sample_size):
```

```
    """
```

```
    Generate data with sin function.
```

```
    :param sample_size: size of data
```

```
    :return: tuple (np.array, np.array)
```

```
    """
```

```
    # generate X
```

```
    X = np.linspace(0, 2 * np.pi, sample_size)
```

```
    # generate y
```

```
    y = np.sin(X) + np.random.normal(loc=0, scale=0.3, size=sample_size)
```

```
    return X.reshape(-1, 1), y.reshape(-1, 1)
```

```
def ridge_regression(feature, target, lam=1e-17):
```

```
    """
```

```
    Compute ridge regression using closed form
```

```
    :param feature: X
```

```
    :param target: y
```

```
    :param lam: lambda
```

```
    :return: parameters (np.array)
```

```
    """
```

```
    feature_dim = feature.shape[1]
```

```
    # TODO: Compute the model of ridge regression.
```

```
    # closed form solution  $w = ((\lambda I + X^T X)^{-1})^T X^T y$ 
```

```
    w = np.linalg.pinv(np.dot(feature.T, feature) + lam *  
np.identity(feature_dim)).dot(feature.T.dot(target))
```

```
    return w
```

```
def k_fold(X, y, k):
```

```
    """
```

```
    Divide data on k samples.
```

```

:param X: np.array(n, m)
:param y: np.array(1, n)
:param k: number of samples to return (int)
:return: list of tuples of 4 np.arrays
"""

# create list for all samples
samples = []
# get indexes and shuffle
indexes = np.array(list(range(len(X))))
# shuffle data
np.random.shuffle(indexes)
# get size of test set
test_size = len(indexes) // k
i = 0
counter = 0
# make folds
while i < len(indexes):
    # get indexes for train and test sets
    if counter == k-1:
        test_idx = indexes[i:]
    else:
        test_idx = indexes[i: i + test_size]
    mask = np.ones(len(indexes), bool)
    mask[test_idx] = False
    train_x, train_y = X[mask], y[mask]
    test_x, test_y = X[test_idx], y[test_idx]
    # add sets to sample
    samples.append((train_x, test_x, train_y, test_y))
    i += test_size
    counter += 1
return samples

```

```

def first_part_assignment():
    # generate sin data
    X, y = generate_sin_data(30)
    # create figure and plot
    plt.figure()
    plt.scatter(X, y)
    yh = np.sin(X)
    arr = sorted(zip(X, yh))
    arr_x, arr_y = list(zip(*arr))
    plt.plot(arr_x, arr_y)
    plt.title('Data points & sin function')
    plt.xlabel('X')
    plt.ylabel('Y')
    plt.show()
    # devide data on train and test sets
    X_train, X_test, y_train, y_test = rand_split_train_test(X, y, 0.3)

```

```

# create list of lambdas
lambdas = [0.0000000001, 0.00001, 0.01, 0.1, 1, 10, 100, 1000]
# create lists for train and test errors
train_errors = []
test_errors = []
# for each lambda
for lmb in lambdas:
    # compute parameters for ridge regression
    w = ridge_regression(X_train, y_train, lmb)
    # make prediction for train set and compute mse
    y_train_predicted = X_train.dot(w)
    train_errors.append(mean_squared_error(y_train, y_train_predicted))
    # make prediction for test set and compute mse
    y_test_predicted = X_test.dot(w)
    test_errors.append(mean_squared_error(y_test, y_test_predicted))
plt.figure()
plt.plot(lambdas, train_errors, label='Train MSE')
plt.plot(lambdas, test_errors, label='Test MSE')
plt.title('MSE for train and test with different lambdas')
plt.xlabel('Lambdas')
plt.ylabel('MSE')
plt.legend()
plt.show()
# devide data on 4 folds
samples = k_fold(X, y, 4)
# create list for errors
errors = []
# for each lambda
for lmb in lambdas:
    mean_test = []
    # for each fold
    # make regression (find parameters, compute test error)
    for sample in samples:
        X_train_sample, X_test_sample, y_train_sample, y_test_sample = sample
        w = ridge_regression(X_train_sample, y_train_sample, lmb)
        y_test_predicted = X_test_sample.dot(w)
        mean_test.append(mean_squared_error(y_test_sample, y_test_predicted))
    # add mean error for every lambda to list
    errors.append(np.mean(mean_test))
# get lambda with minimum test error
print('The best choice of lambda: {}'.format(lambdas[np.argmin(errors)]))

```

2. (30 points) Complete the gradient descent optimizers (gradient descent and stochastic gradient descent). Use the code to conduct the following experiments:

- (a) (10 points) Implement ridge regression using the gradient descent optimizer. Apply gradient descent based ridge regression using a random data set with $N = 1000$ and $d = 50$, and plot the objective changes at each iteration.
- (b) (20 points) Apply stochastic gradient descent using a batch size $n = 5; 10; 100; 500$ and for each batch size, plot the objective changes at each iteration.

Solution

```
def get_batches(X, Y, size):  
    """  
    Divide data into list of mini-batches.  
  
    :param X: np.array(n, m)  
    :param y: np.array(1, n)  
    :param size: size of mini-batch (int)  
    :return: list of tuples (np.array(n, size), np.array(1, size))  
    """  
  
    # get length of X (and y) - the same  
    n_samples = X.shape[0]  
    # shuffle data  
    indexes = np.array(list(range(X.shape[0])))  
    np.random.shuffle(indexes)  
    x = X[indexes]  
    y = Y[indexes]  
    # create list for mini-batches  
    samples = []  
    # for i from 0 to length of dataset with step of size of mini-batch  
    for i in range(0, n_samples, size):  
        # find begin and end index of future mini-batch  
        begin, end = int(i), int(min(i + size, n_samples))  
        # get X and y mini-batches and add them into list  
        samples.append((x[begin:end, :], y[begin:end]))  
    return samples  
  
def ridge_regression_GD(X, y, lmbda=0.01, l_rate=0.1, tol=0.000001):  
    """  
    Ridge stochastic gradient descent.  
  
    :param X: np.array(n, m)  
    :param y: np.array(1, n)  
    :param lmbda: parameter for penalty of 2 (float or int)  
    :param l_rate: step of GD (float or int)  
    :param tol: tolerance (criteria for stopping)  
    :return: tuple(np.array(X.shape[0], 1), list of floats)  
    """  
  
    # initialize w: parameters of the model  
    w = np.zeros((X.shape[1], 1))  
    # create list for losses  
    errors = []  
    converged = False  
    i = 1  
    old_error = 0  
    # while algorithm does not converge  
    while not converged:  
        # get prediction
```

```

yhat = X.dot(w)
# start to compute gradient
# get error
error = yhat - y
# update parameters
w -= l_rate * (X.T.dot(error) + lmbda * w) / X.shape[0]
# get loss
new_error = (np.sum((y - X.dot(w))**2) + lmbda * np.linalg.norm(w)) / (2 * X.shape[0])
# check how loss from last and current iteration changes
# if difference between current and previous loss is less than tolerance
# algorithm converged
if np.abs(old_error - new_error) < tol:
    converged = True
# increase number of iteration
i += 1
# save current loss
errors.append(new_error)
# update old loss to current
old_error = new_error
# print loss if iteration number divides by 100
if i % 100 == 0:
    print(i, new_error)
return w, errors

```

```

def ridge_regression_SGD(X, Y, batch, lmbda=0.01, l_rate=0.01, tol=0.0000001):

```

```

    """

```

```

    Ridge regression stochastic gradient descent.

```

```

:param X: np.array(n, m)
:param y: np.array(1, n)
:param batch: size of batch (int)
:param lmbda: parameter for penalty of 2 (float or int)
:param l_rate: step of GD (float or int)
:param tol: tolerance (criteria for stopping)
:return: tuple(np.array(X.shape[0], 1), list of floats)
    """

```

```

# initialize w: parameters of the model

```

```

w = np.zeros((X.shape[1], 1))

```

```

# create list for losses

```

```

errors = []

```

```

# divide data on mini-batches

```

```

batches = get_batches(X, Y, batch)

```

```

converged = False

```

```

i = 1

```

```

old_error = 0

```

```

# while algorithm does not converge

```

```

while not converged:

```

```

    # create list for store loss on each batch

```

```

    batch_errors = []

```

```

    # for every part of X and y

```

```

for (x, y) in batches:
    # get prediction
    yhat = x.dot(w)
    # get error
    error = yhat - y
    # update parameter
    w -= l_rate * (x.T.dot(error) + lmbda * w) / X.shape[0]
    # compute loss on current batch
    batch_error = (np.sum((y - x.dot(w))**2) + lmbda * np.linalg.norm(w)) / (2 *
X.shape[0] * i)
    batch_errors.append(batch_error)
    # add to list of all losses
    errors.append(batch_error)
    # compute mean loss of current iteration
    new_error = np.mean(batch_errors)
    # check how loss from last and current iteration changes
    # if difference between current and previous loss is less than tolerance
    # algorithm converged
    if np.abs(old_error - new_error) < tol:
        converged = True
    # increase number of current iteration
    i += 1
    # update current loss
    old_error = new_error
    if i % 100 == 0:
        print(i, new_error)
return w, errors

```

```

def second_part_assignment():
    # generate random data with 1000 samples and 30 features
    X = np.random.randn(50).reshape(1, -1)
    for i in range(999):
        X = np.vstack((X, np.random.randn(50).reshape(1, -1)))
    y = np.random.randn(1000) * np.random.randn(1000)
    y = y.reshape(-1, 1)
    # compute model parameters and losses
    w, errors = ridge_regression_GD(X, y, l_rate=0.01, lmbda=0.01, tol=0.00000001)
    # plot results
    plt.figure()
    plt.plot(list(range(len(errors))), errors)
    plt.xlabel('Number of iterations')
    plt.ylabel('Loss')
    plt.title('Gradient descent ridge regression loss')
    plt.show()
    print('Got MSE: {}'.format(mean_squared_error(y, X.dot(w))))
    # create list for errors
    mean_errors = []
    # for different sizes of batches
    for batch in [5, 10, 100, 500]:
        # create plot

```

```

plt.figure()
# compute parameters and losses with stochastic GD
w, errors = ridge_regression_SGD(X, y, lmbda=0.1, l_rate=0.1, batch=batch)
# compute mse
mean_errors.append(mean_squared_error(y, X.dot(w)))
# plot losses
plt.plot(list(range(len(errors))), errors)
plt.title('SGD Ridge regression loss with batch size {}'.format(batch))
plt.xlabel('Number of iterations')
plt.ylabel('Loss')
plt.show()

# create plot for mse of ridge with parameters got while training SGD
plt.figure()
plt.plot([5, 10, 100, 500], mean_errors)
plt.title('MSE')
plt.xlabel('Batch size')
plt.ylabel('Loss')
plt.show()

```

Whole solution in regression_class.py

```

import time
import numpy as np
from sklearn.utils import shuffle
import matplotlib.pyplot as plt

```

```

def rand_split_train_test(data, label, train_perc):
    if train_perc >= 1 or train_perc <= 0:
        raise Exception('train_perc should be between (0,1).')
    sample_size = data.shape[0]
    if sample_size < 2:
        raise Exception('Sample size should be larger than 1. ')

    num_train_sample = np.max([np.floor(sample_size * train_perc).astype(int), 1])
    data, label = shuffle(data, label)

    data_tr = data[:num_train_sample]
    data_te = data[num_train_sample:]

    label_tr = label[:num_train_sample]
    label_te = label[num_train_sample:]

    return data_tr, data_te, label_tr, label_te

```

```

def subsample_data(data, label, subsample_size):
    # protected sample size
    subsample_size = np.max([1, np.min([data.shape[0], subsample_size])])

```

```

data, label = shuffle(data, label)
data = data[:subsample_size]
label = label[:subsample_size]
return data, label

```

```

def generate_rnd_data(feature_size, sample_size, bias=False):

```

```

    # Generate X matrix.
    data = np.concatenate((np.random.randn(sample_size, feature_size), np.ones((sample_size,
1))), axis=1) \
        if bias else np.random.randn(sample_size, feature_size) # the first dimension is sample_size
    (n X d)

    # Generate ground truth model.
    truth_model = np.random.randn(feature_size + 1, 1) * 10 \
        if bias else np.random.randn(feature_size, 1) * 10

    # Generate label.
    label = np.dot(data, truth_model)

    # add element-wise gaussian noise to each label.
    label += np.random.randn(sample_size, 1)
    return data, label, truth_model

```

```

def mean_squared_error(true_label, predicted_label):

```

```

    """
    Compute the mean square error between the true and predicted labels
    :param true_label: Nx1 vector
    :param predicted_label: Nx1 vector
    :return: scalar MSE value
    """
    mse = np.sqrt(np.sum((true_label - predicted_label)**2)/true_label.size)
    return mse

```

```

def least_squares(feature, target):

```

```

    """
    Compute least squares using closed form
    :param feature: X
    :param target: y
    :return: computed weight vector
    """

    # TODO: Compute the model of least squares.
    w = np.linalg.pinv(np.dot(feature.T, feature)).dot(feature.T.dot(target))
    return w

```



```
def generate_sin_data(sample_size):
    """
    Generate data with sin function.

    :param sample_size: size of data
    :return: tuple (np.array, np.array)
    """
    # generate X
    X = np.linspace(0, 2 * np.pi, sample_size)
    # generate y
    y = np.sin(X) + np.random.normal(loc=0, scale=0.3, size=sample_size)
    return X.reshape(-1, 1), y.reshape(-1, 1)
```

```
def ridge_regression(feature, target, lam=1e-17):
    """
    Compute ridge regression using closed form

    :param feature: X
    :param target: y
    :param lam: lambda
    :return: parameters (np.array)
    """
    feature_dim = feature.shape[1]

    # TODO: Compute the model of ridge regression.
    # closed form solution  $w = ((\lambda I + X.T @ X)^{-1}) @ X.T @ y$ 
    w = np.linalg.pinv(np.dot(feature.T, feature) + lam *
np.identity(feature_dim)).dot(feature.T.dot(target))
    return w
```

```
def get_batches(X, Y, size):
    """
    Divide data into list of mini-batches.

    :param X: np.array(n, m)
    :param y: np.array(1, n)
    :param size: size of mini-batch (int)
    :return: list of tuples (np.array(n, size), np.array(1, size))
    """
    # get length of X (and y) - the same
    n_samples = X.shape[0]
    # shuffle data
    indexes = np.array(list(range(X.shape[0])))
    np.random.shuffle(indexes)
    x = X[indexes]
    y = Y[indexes]
    # create list for mini-batches
    samples = []
    # for i from 0 to length of dataset with step of size of mini-batch
    for i in range(0, n_samples, size):
```

```

    # find begin and end index of future mini-batch
    begin, end = int(i), int(min(i + size, n_samples))
    # get X and y mini-batches and add them into list
    samples.append((x[begin:end, :], y[begin:end]))
return samples

```

```

def ridge_regression_GD(X, y, lmbda=0.01, l_rate=0.1, tol=0.000001):

```

```

    """
    Ridge stochastic gradient descent.

    :param X: np.array(n, m)
    :param y: np.array(1, n)
    :param lmbda: parameter for penalty of 2 (float or int)
    :param l_rate: step of GD (float or int)
    :param tol: tolerance (criteria for stopping)
    :return: tuple(np.array(X.shape[0], 1), list of floats)
    """

    # initialize w: parameters of the model
    w = np.zeros((X.shape[1], 1))
    # create list for losses
    errors = []
    converged = False
    i = 1
    old_error = 0
    # while algorithm does not converge
    while not converged:
        # get prediction
        yhat = X.dot(w)
        # start to compute gradient
        # get error
        error = yhat - y
        # update parameters
        w -= l_rate * (X.T.dot(error) + lmbda * w) / X.shape[0]
        # get loss
        new_error = (np.sum((y - X.dot(w))**2) + lmbda * np.linalg.norm(w)) / (2 * X.shape[0])
        # check how loss from last and current iteration changes
        # if difference between current and previous loss is less than tolerance
        # algorithm converged
        if np.abs(old_error - new_error) < tol:
            converged = True
        # increase number of iteration
        i += 1
        # save current loss
        errors.append(new_error)
        # update old loss to current
        old_error = new_error
        # print loss if iteration number divides by 100
        if i % 100 == 0:

```

```
    print(i, new_error)
    return w, errors
```

```
def ridge_regression_SGD(X, Y, batch, lmbda=0.01, l_rate=0.01, tol=0.0000001):
```

```
    """
```

```
    Ridge regression stochastic gradient descent.
```

```
    :param X: np.array(n, m)
    :param y: np.array(1, n)
    :param batch: size of batch (int)
    :param lmbda: parameter for penalty of 2 (float or int)
    :param l_rate: step of GD (float or int)
    :param tol: tolerance (criteria for stopping)
    :return: tuple(np.array(X.shape[0], 1), list of floats)
    """
```

```
    # initialize w: parameters of the model
```

```
    w = np.zeros((X.shape[1], 1))
```

```
    # create list for losses
```

```
    errors = []
```

```
    # divide data on mini-batches
```

```
    batches = get_batches(X, Y, batch)
```

```
    converged = False
```

```
    i = 1
```

```
    old_error = 0
```

```
    # while algorithm does not converge
```

```
    while not converged:
```

```
        # create list for store loss on each batch
```

```
        batch_errors = []
```

```
        # for every part of X and y
```

```
        for (x, y) in batches:
```

```
            # get prediction
```

```
            yhat = x.dot(w)
```

```
            # get error
```

```
            error = yhat - y
```

```
            # update parameter
```

```
            w -= l_rate * (x.T.dot(error) + lmbda * w) / X.shape[0]
```

```
            # compute loss on current batch
```

```
            batch_error = (np.sum((y - x.dot(w))**2) + lmbda * np.linalg.norm(w)) / (2 * X.shape[0] * i)
```

```
            batch_errors.append(batch_error)
```

```
            # add to list of all losses
```

```
            errors.append(batch_error)
```

```
        # compute mean loss of current iteration
```

```
        new_error = np.mean(batch_errors)
```

```
        # check how loss from last and current iteration changes
```

```
        # if difference between current and previous loss is less than tolerance
```

```
        # algorithm converged
```

```
        if np.abs(old_error - new_error) < tol:
```

```
            converged = True
```

```
        # increase number of current iteration
```

```
        i += 1
```

```

    # update current loss
    old_error = new_error
    if i % 100 == 0:
        print(i, new_error)
    return w, errors

```

```

def k_fold(X, y, k):

```

```

    """

```

```

    Divide data on k samples.

```

```

    :param X: np.array(n, m)

```

```

    :param y: np.array(1, n)

```

```

    :param k: number of samples to return (int)

```

```

    :return: list of tuples of 4 np.arrays

```

```

    """

```

```

    # create list for all samples

```

```

    samples = []

```

```

    # get indexes and shuffle

```

```

    indexes = np.array(list(range(len(X))))

```

```

    # shuffle data

```

```

    np.random.shuffle(indexes)

```

```

    # get size of test set

```

```

    test_size = len(indexes) // k

```

```

    i = 0

```

```

    counter = 0

```

```

    # make folds

```

```

    while i < len(indexes):

```

```

        # get indexes for train and test sets

```

```

        if counter == k-1:

```

```

            test_idx = indexes[i:]

```

```

        else:

```

```

            test_idx = indexes[i: i + test_size]

```

```

            mask = np.ones(len(indexes), bool)

```

```

            mask[test_idx] = False

```

```

            train_x, train_y = X[mask], y[mask]

```

```

            test_x, test_y = X[test_idx], y[test_idx]

```

```

            # add sets to sample

```

```

            samples.append((train_x, test_x, train_y, test_y))

```

```

            i += test_size

```

```

            counter += 1

```

```

    return samples

```

```

def exp1():

```

```

    # EXP1: training testing.

```

```

    # generate a data set.

```

```

    (feature_all, target_all, model) = generate_rnd_data(feature_size=3, sample_size=20,

```

```

bias=True)
    # split training/testing
    feature_train, feature_test, target_train, target_test = rand_split_train_test(feature_all,
target_all, train_perc=0.8)
    # compute model
    reg_model_lsqr = least_squares(feature_train, target_train)
    reg_model_ridge = ridge_regression(feature_train, target_train, lam=1e-7)

    # evaluate performance
    print('Training MSE(lsqr):', mean_squared_error(target_train, np.dot(feature_train,
reg_model_lsqr)))
    print('Testing MSE(lsqr):', mean_squared_error(target_test, np.dot(feature_test,
reg_model_lsqr)))
    print('Training MSE(ridge):', mean_squared_error(target_train, np.dot(feature_train,
reg_model_ridge)))
    print('Testing MSE(ridge):', mean_squared_error(target_test, np.dot(feature_test,
reg_model_ridge)))

def exp2():
    # EXP2: generalization performance: increase sample size.
    different_sample_sizes = [50, 100, 150, 200, 250, 300, 350, 400, 450]
    (feature_all, target_all, model) = generate_rnd_data(feature_size=100, sample_size=1000,
bias=True)
    feature_hold, feature_test, target_hold, target_test = \
        rand_split_train_test(feature_all, target_all, train_perc=0.9)

    train_performance = []
    test_performance = []
    for train_sample_size in different_sample_sizes:
        feature_train, target_train = subsample_data(feature_hold, target_hold, train_sample_size)
        reg_model = ridge_regression(feature_train, target_train, lam=1e-5)
        train_performance += [mean_squared_error(target_train, np.dot(feature_train, reg_model))]
        test_performance += [mean_squared_error(target_test, np.dot(feature_test, reg_model))]

    print(train_performance)
    print(test_performance)

    plt.figure()
    train_plot, = plt.plot(different_sample_sizes, np.log10(train_performance), linestyle='-',
color='b',
                        label='Training Error')
    test_plot, = plt.plot(different_sample_sizes, np.log10(test_performance), linestyle='-', color='r',
label='Testing '
                        'Error')

    plt.xlabel("Sample Size")
    plt.ylabel("Error (log)")
    plt.title("Generalization performance: increase sample size fix dimensionality")
    plt.legend(handles=[train_plot, test_plot])
    plt.show()

```

def exp3():

EXP3: generalization performance: increase dimensionality.

different_dimensionality = [100, 150, 200, 250, 300, 350, 400, 450]

train_performance = []

test_performance = []

for dimension **in** different_dimensionality:

(feature_all, target_all, model) = generate_rnd_data(feature_size=dimension,
sample_size=1000, bias=True)

feature_train, feature_test, target_train, target_test = \

rand_split_train_test(feature_all, target_all, train_perc=0.9)

reg_model = ridge_regression(feature_train, target_train, lam=1e-5)

train_performance += [mean_squared_error(target_train, np.dot(feature_train, reg_model))]

test_performance += [mean_squared_error(target_test, np.dot(feature_test, reg_model))]

print(train_performance)

print(test_performance)

plt.figure()

train_plot, = plt.plot(different_dimensionality, np.log10(train_performance), linestyle='-',
color='b',

label='Training Error')

test_plot, = plt.plot(different_dimensionality, np.log10(test_performance), linestyle='-',
color='r', label='Testing ')

'Error')

plt.xlabel("Dimensionality")

plt.ylabel("Error (log)")

plt.title("Generalization performance: increase dimensionality fix sample size")

plt.legend(handles=[train_plot, test_plot])

plt.show()

def exp4():

EXP4: computational time: increase dimensionality.

different_dimensionality = range(100, 2000, 100)

train_performance = []

test_performance = []

time_elapse = []

for dimension **in** different_dimensionality:

(feature_all, target_all, model) = generate_rnd_data(feature_size=dimension,
sample_size=1000, bias=True)

feature_train, feature_test, target_train, target_test = \

rand_split_train_test(feature_all, target_all, train_perc=0.9)

t = time.time()

reg_model = ridge_regression(feature_train, target_train, lam=1e-5)

time_elapse += [time.time() - t]

print('Finished model of dimension {}'.format(dimension))

```

train_performance += [mean_squared_error(target_train, np.dot(feature_train, reg_model))]
test_performance += [mean_squared_error(target_test, np.dot(feature_test, reg_model))]

plt.figure()
time_plot, = plt.plot(different_dimensionality, time_elapse, linestyle='-', color='r', label='Time
cost')
plt.xlabel("Dimensionality")
plt.ylabel("Time (ms)")
plt.title("Computational efficiency.")
plt.legend(handles=[time_plot])
plt.show()

```

```

def first_part_assignment():
    # generate sin data
    X, y = generate_sin_data(30)
    # create figure and plot
    plt.figure()
    plt.scatter(X, y)
    yh = np.sin(X)
    arr = sorted(zip(X, yh))
    arr_x, arr_y = list(zip(*arr))
    plt.plot(arr_x, arr_y)
    plt.title('Data points & sin function')
    plt.xlabel('X')
    plt.ylabel('Y')
    plt.show()
    # divide data on train and test sets
    X_train, X_test, y_train, y_test = rand_split_train_test(X, y, 0.3)
    # create list of lambdas
    lambdas = [0.0000000001, 0.00001, 0.01, 0.1, 1, 10, 100, 1000]
    # create lists for train and test errors
    train_errors = []
    test_errors = []
    # for each lambda
    for lmb in lambdas:
        # compute parameters for ridge regression
        w = ridge_regression(X_train, y_train, lmb)
        # make prediction for train set and compute mse
        y_train_predicted = X_train.dot(w)
        train_errors.append(mean_squared_error(y_train, y_train_predicted))
        # make prediction for test set and compute mse
        y_test_predicted = X_test.dot(w)
        test_errors.append(mean_squared_error(y_test, y_test_predicted))
    plt.figure()
    plt.plot(lambdas, train_errors, label='Train MSE')
    plt.plot(lambdas, test_errors, label='Test MSE')
    plt.title('MSE for train and test with different lambdas')
    plt.xlabel('Lambdas')

```

```

plt.ylabel('MSE')
plt.legend()
plt.show()
# divide data on 4 folds
samples = k_fold(X, y, 4)
# create list for errors
errors = []
# for each lambda
for lmb in lambdas:
    mean_test = []
    # for each fold
    # make regression (find parameters, compute test error)
    for sample in samples:
        X_train_sample, X_test_sample, y_train_sample, y_test_sample = sample
        w = ridge_regression(X_train_sample, y_train_sample, lmb)
        y_test_predicted = X_test_sample.dot(w)
        mean_test.append(mean_squared_error(y_test_sample, y_test_predicted))
    # add mean error for every lambda to list
    errors.append(np.mean(mean_test))
# get lambda with minimum test error
print('The best choice of lambda: {}'.format(lambdas[np.argmin(errors)]))

```

```

def second_part_assignment():
    # generate random data with 1000 samples and 30 features
    X = np.random.randn(50).reshape(1, -1)
    for i in range(999):
        X = np.vstack((X, np.random.randn(50).reshape(1, -1)))
    y = np.random.randn(1000) * np.random.randn(1000)
    y = y.reshape(-1, 1)
    # compute model parameters and losses
    w, errors = ridge_regression_GD(X, y, l_rate=0.01, lmbda=0.01, tol=0.00000001)
    # plot results
    plt.figure()
    plt.plot(list(range(len(errors))), errors)
    plt.xlabel('Number of iterations')
    plt.ylabel('Loss')
    plt.title('Gradient descent ridge regression loss')
    plt.show()
    print('Got MSE: {}'.format(mean_squared_error(y, X.dot(w))))
    # create list for errors
    mean_errors = []
    # for different sizes of batches
    for batch in [5, 10, 100, 500]:
        # create plot
        plt.figure()
        # compute parameters and losses with stochastic GD
        w, errors = ridge_regression_SGD(X, y, lmbda=0.1, l_rate=0.1, batch=batch)
        # compute mse
        mean_errors.append(mean_squared_error(y, X.dot(w)))

```



```

    # plot losses
    plt.plot(list(range(len(errors))), errors)
    plt.title('SGD Ridge regression loss with batch size {}'.format(batch))
    plt.xlabel('Number of iterations')
    plt.ylabel('Loss')
    plt.show()

    # create plot for mse of ridge with parameters got while training SGD
    plt.figure()
    plt.plot([5, 10, 100, 500], mean_errors)
    plt.title('MSE')
    plt.xlabel('Batch size')
    plt.ylabel('Loss')
    plt.show()

if __name__ == '__main__':
    plt.interactive(False)

    # set seeds to get repeatable results.
    np.random.seed(491)

    # # EXP1: training testing.
    # exp1()
    #
    # # EXP2: generalization performance: increase sample size.
    # exp2()
    #
    # # EXP3: generalization performance: increase dimensionality.
    # exp3()
    #
    # # EXP4: computational complexity by varing dimensions.
    # exp4()
    first_part_assignment()
    second_part_assignment()

```